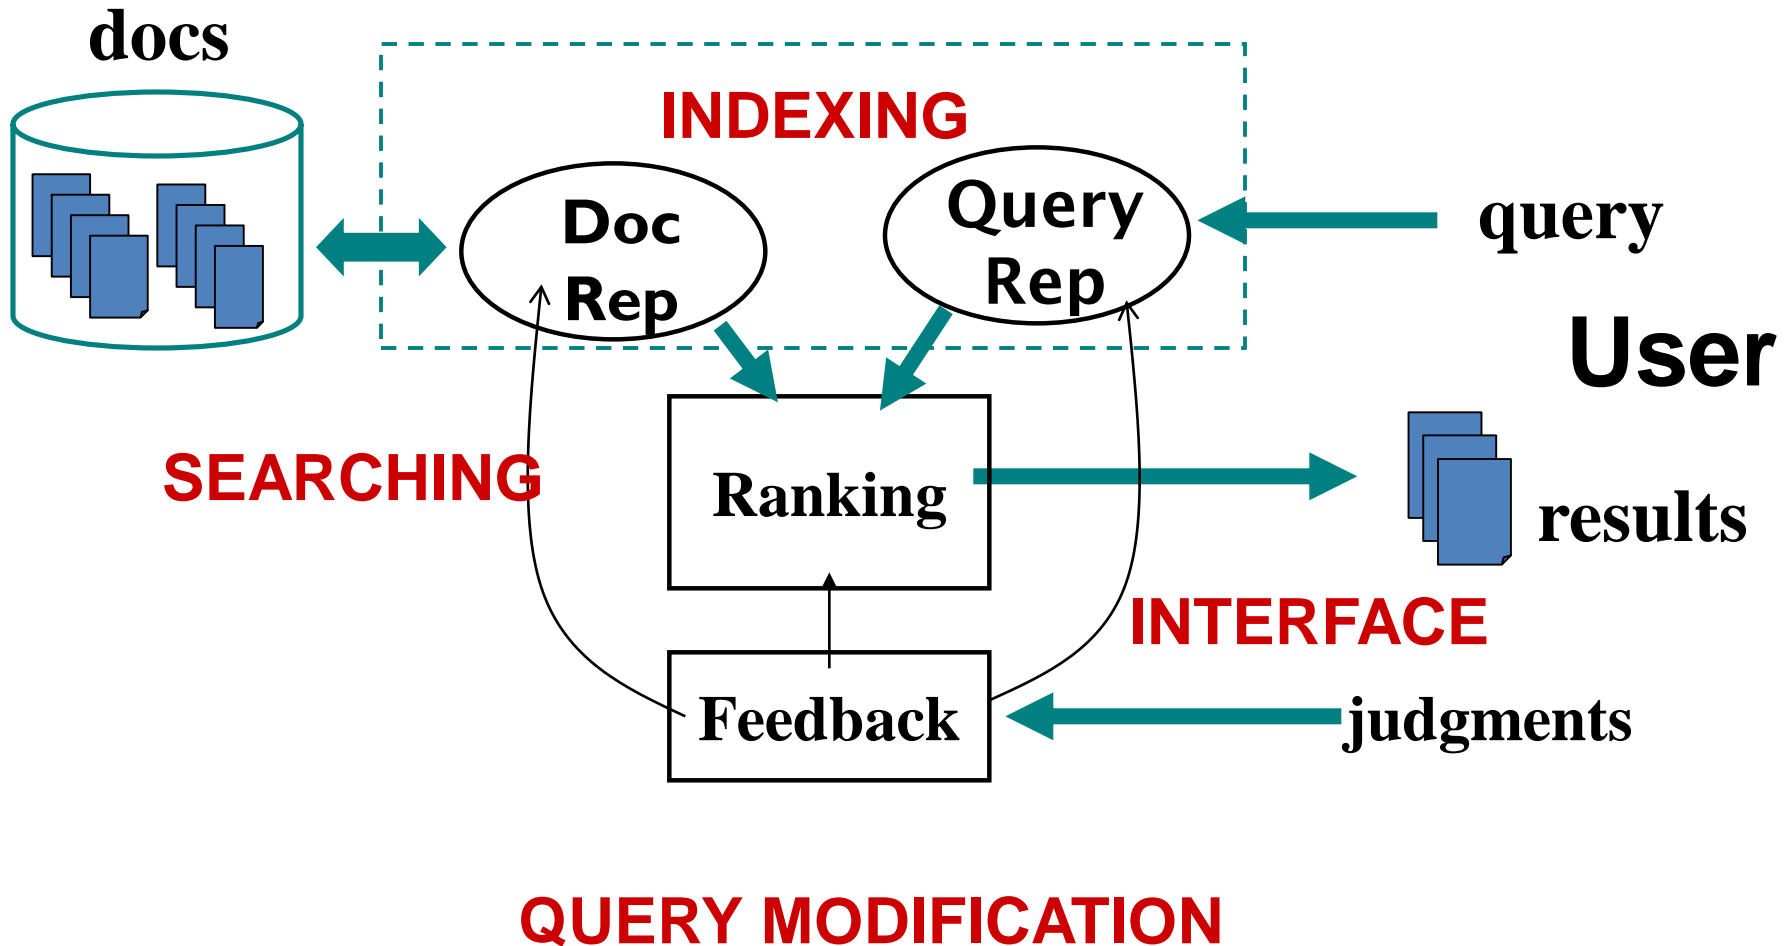


Implementation Issues & IR Systems

Lecture Plan

- How to implement a simple IR system
 - Index construction
 - Scoring
- Open source IR toolkits

IR System Architecture



Indexing

- Indexing = Convert documents to data structures that enable fast search

Unstructured data

- Which plays of Shakespeare contain the words Brutus and Caesar, but not Calpurnia?
- One could grep all of Shakespeare's plays for Brutus and Caesar, then strip out lines containing Calpurnia.
- Why is grep not the solution?
 - Slow (for large collections)
 - “not Calpurnia” is non-trivial
 - Other operations (e.g., find the word Romans near countryman) not feasible
 - Ranked retrieval (best documents to return)

Term-document incidence matrix

	Anthony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Mac beth	...
Anthony	1	1	0	0	0	1	
Brutus	1	1	0	1	0	0	
Caesar	1	1	0	1	1	1	
Calpurnia	0	1	0	0	0	0	
...							

Entry is 1 if term occurs. Example: Calpurnia occurs in Julius Caesar.
Entry is 0 if term doesn't occur. Example: Calpurnia doesn't occur in The tempest.

Incidence vectors

- So we have a 0/1 vector for each term.
- To answer the query Brutus and Caesar and not Calpurnia:
 - Take the vectors for Brutus, Caesar, and Calpurnia
 - Complement the vector of Calpurnia
 - Do a (bitwise) **AND** on the three vectors
 - **110100 AND 110111 AND 101111 = 100100**

Bigger collections

- Consider $N = 10^6$ documents, each with about 1000 tokens
- On average 6 bytes per token, including spaces and punctuation \Rightarrow size of document collection is about 6 GB
- Assume there are $M = 500000$ distinct terms in the collection
- $M = 500,000 \times 10^6 =$ half a trillion 0s and 1s.
- But the matrix has no more than one billion 1s.
 - Matrix is extremely sparse.
- What is a better representations?
 - We only record the 1s.

Indexing

- Inverted index is the dominating indexing method (used by all search engines)
- Other indices (e.g., document index) may be needed for feedback

Inverted Index

- Fast access to all docs containing a given term (along with freq and pos information)
- For each term, we get a list of tuples (docID, freq, pos).
- Given a query, we can fetch the lists for all query terms and work on the involved documents.
 - Boolean query: set operation
 - Natural language query: term weight summing

Inverted Index Example

Doc 1

... news about

Doc 2

... news
about
organic food
campaign

Doc 3

... news of presidential campaign
... presidential candidate ...

Dictionary (or Lexicon)

Term	# docs	Total freq
news	3	3
campaign	2	2
presidential	1	2
food	1	1
...

Postings

Doc id	Freq
1	1
2	1
3	1
2	1
3	1
2	2
2	1
...	...
...	...

Position

p1

p2

p3

p4

p5

p6, p7


p8

Inverted Index for Fast Search?

- Single-term query?
- Multi-term Boolean query?
 - Must match term “A” AND term “B”
 - Must match term “A” OR term “B”
- Multi-term keyword query
 - Similar to disjunctive Boolean query (“A” OR “B”)
 - Aggregate term weights
- More efficient than sequentially scanning documents (why?)

Empirical Distribution of Words – Zipf's Law

- $rank \times frequency$ constant



Data Structures for Inverted Index

- Dictionary: modest size
 - Needs fast random access
 - Preferred to be in memory
 - Hash table, B-tree, trie, ...
- Postings: huge
 - Sequential access is expected
 - Can stay on disk
 - May contain docID, term freq., term pos, etc
 - Compression is desirable

Inverted Index Compression

- In general, leverage skewed distribution of values and use variable-length encoding
- TF compression
 - Small numbers tend to occur far more frequently than large numbers (why?)
 - Fewer bits for small (high frequency) integers at the cost of more bits for large integers
- Doc ID compression
 - “d-gap” (store difference): $d_1, d_2 - d_1, d_3 - d_2, \dots$
 - Feasible due to sequential access
- Methods: Binary code, unary code, γ -code, δ -code, ...

Integer Compression Methods

- Binary: equal-length coding
- Unary: $x \geq 1$ is coded as x one bits followed by 0, e.g., $3 \Rightarrow 1110$; $5 \Rightarrow 111110$
- γ -code: $x \Rightarrow$ a pair of length, offset. *Offset* is x in binary with the leading 1 removed. *Length* encodes the length of offset in unary code
- δ -code: same as γ -code, but replace the unary prefix with γ -code.

Integer Compression Methods - Example

X = 23

- Unary: 1111111111111111111110
- γ -code: $x \Rightarrow$ a pair of length, offset. *Offset* is x in binary with the leading 1 removed. *Length* encodes the length of offset in unary code
 - 23 in binary: 10111
 - After removing leading 1: 0111 (offset)
 - Length of offset in unary: 11110 (length)
 - 111100111

Integer Compression Methods - Example

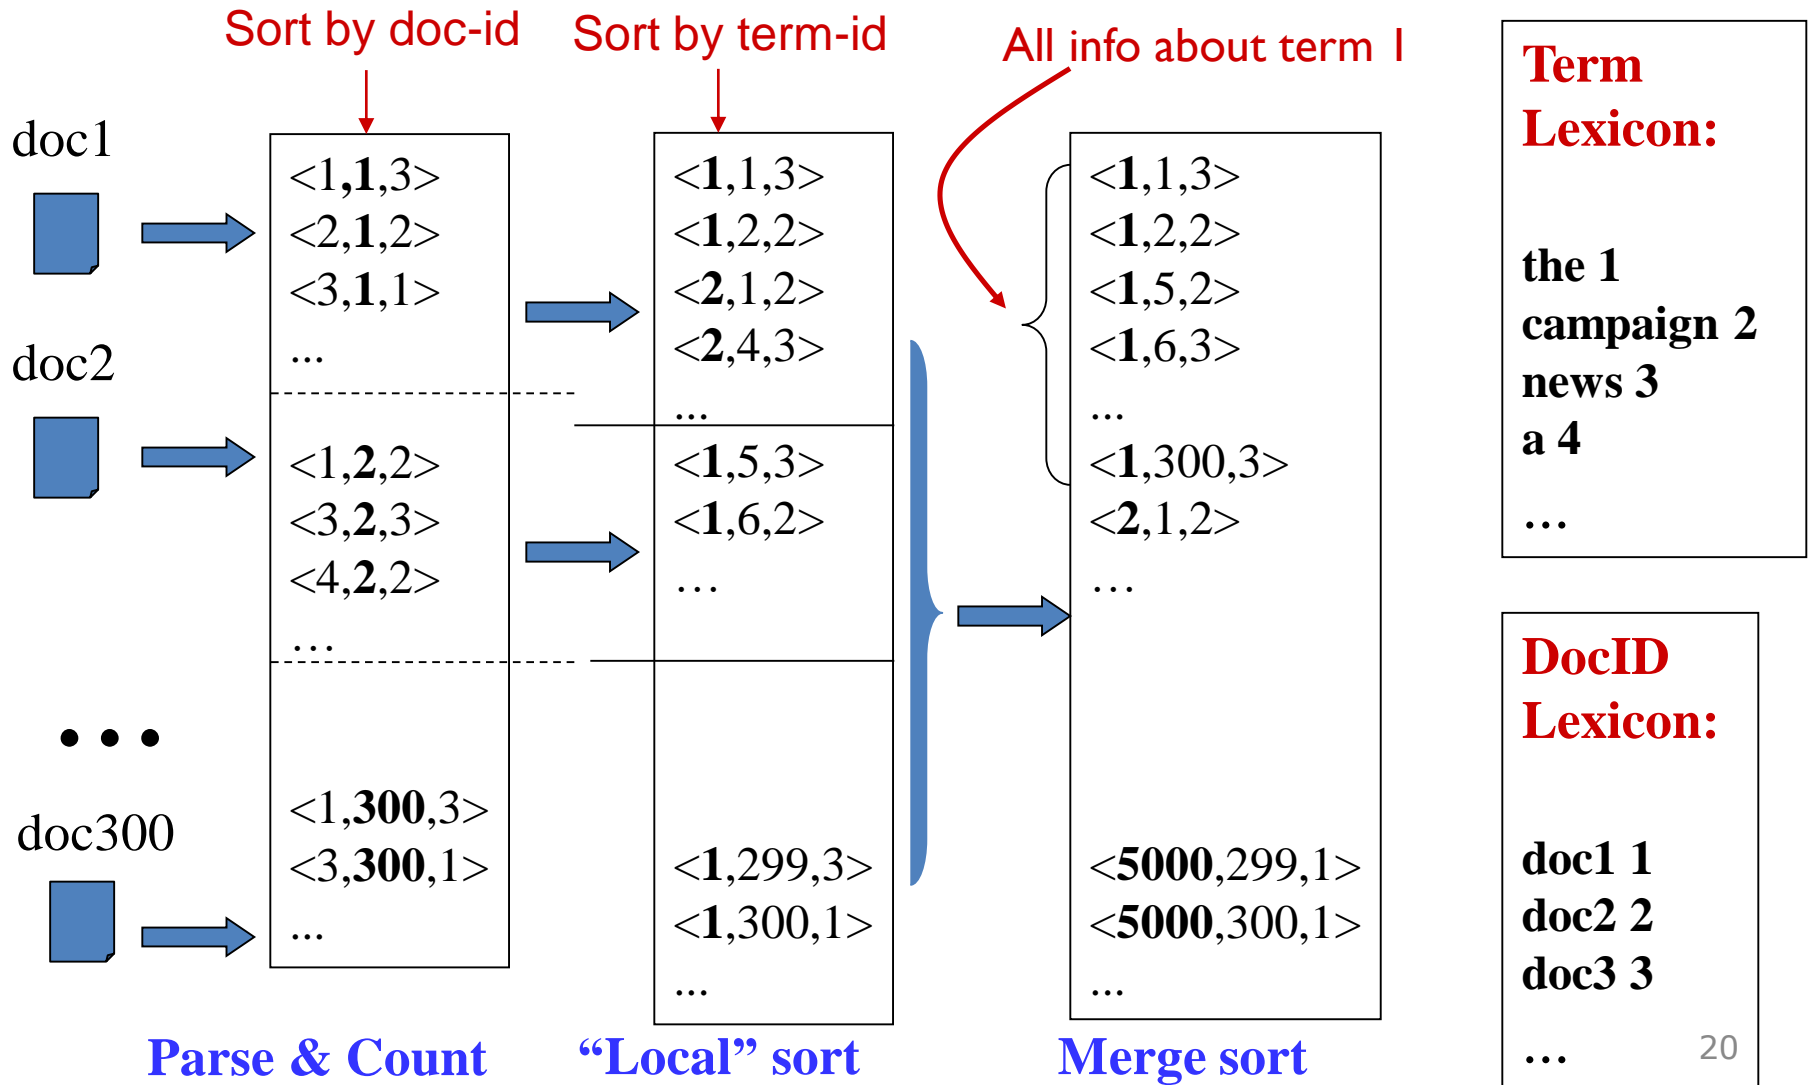
$$\underline{X = 23}$$

- δ -code: same as γ -code ,but replace the unary prefix with γ -code.
 - Offset: 0111
 - Length in γ -code: 11000
 - 110000111

Constructing Inverted Index

- The main difficulty is to build a huge index with limited memory
- Memory-based methods: not usable for large collections
- Sort-based methods:
 - Step 1: collect local (termID, docID, freq) tuples
 - Step 2: sort local tuples (to make “runs”)
 - Step 3: pair-wise merge runs
 - Step 4: Output inverted file

Sort-based Inversion



Searching

- Given a query, score documents efficiently
- Boolean query
 - Fetch the inverted list for all query terms
 - Perform set operations to get the subset of docs that satisfy the Boolean condition
 - E.g., $Q_1 = \text{"info" AND "security"}$, $Q_2 = \text{"info" OR "security"}$
 - info: d_1, d_2, d_3, d_4
 - security: d_2, d_4, d_6
 - Results: $\{d_2, d_4\} (Q_1) \{d_1, d_2, d_3, d_4, d_6\} (Q_2)$

How to Score Documents Quickly?

General form of scoring functions

The diagram illustrates the general form of scoring functions $f(q, d)$. The equation is $f(q, d) = f_a(h(g(t_1, d, q), \dots, g(t_k, d, q)), f_d(d), f_q(q))$. Annotations include: 'Final score adjustment' pointing to f_a ; 'Weight aggregation' pointing to h ; 'Weight of a matched query term in d' pointing to g ; and arrows pointing from $f_d(d)$ and $f_q(q)$ to f_a .

$$f(q, d) = f_a(h(g(t_1, d, q), \dots, g(t_k, d, q)), f_d(d), f_q(q))$$

Final score adjustment

Weight **aggregation**

Weight of a **matched** query term in d

A General Algorithm for Ranking Documents

$$f(q, d) = f_a \left(h(g(t_1, d, q), \dots, g(t_k, d, q)), f_d(d), f_q(q) \right)$$

- $f_d(d)$ and $f_q(q)$ are pre-computed
- Maintain a score accumulator for each d to compute h
- For each query term t_i
 - Fetch the inverted list $\{(d_1, f_1), \dots, (d_n, f_n)\}$
 - For each entry (d_j, f_j) , compute $g(t_i, d_j, q)$, and update score accumulator for d_j to incrementally compute h
- Adjust the score to compute f_a , and sort

Ranking Documents: Example

$f(d, q) = g(t_1, d, q) + \dots + g(t_k, d, q)$ where $g(t_i, d, q) = c(t_i, d)$

Query = “info security”

Info: $(d_1, 3), (d_2, 4), (d_3, 1), (d_4, 5)$

Security: $(d_2, 3), (d_4, 1), (d_5, 3)$

Accumulators:		d_1	d_2	d_3	d_4	d_5
		0	0	0	0	0
info	$(d_1, 3) \Rightarrow$	3	0	0	0	0
	$(d_2, 4) \Rightarrow$	3	4	0	0	0
	$(d_3, 1) \Rightarrow$	3	4	1	0	0
	$(d_4, 5) \Rightarrow$	3	4	1	5	0
security	$(d_2, 3) \Rightarrow$	3	7	1	5	0
	$(d_4, 1) \Rightarrow$	3	7	1	6	0
	$(d_5, 3) \Rightarrow$	3	7	1	6	3

Further Improving Efficiency

- Caching (e.g., query results, list of inverted index)
- Keep only the most promising accumulators
- Sort the inverted list in decreasing order of weights and fetch only N entries with the highest weights
- Scaling up to the Web-scale (more about this later)

Some Text Retrieval Toolkits

- Smart (Cornell) (**no longer popular**)
- Lucene (<http://lucene.apache.org/>)
- Lemur/Indri (<http://lemurproject.org/>)
- Galago
- Terrier (<http://terrier.org/>)
- MeTA (<http://meta-toolkit.github.io/meta/>)

Questions?